



# Autoencoders and Deep Representation Learning



Center for  
Advanced  
Studies in  
Adaptive  
Systems

by Alireza Ghods



# Outline

- Quick review
- Background
- Autoencoder
- Autoencoder + regularizer
- Variational Autoencoder
- Questions



# Review: Feature Engineering vs. Representation Learning



Feature  
engineering

Edges?  
Texture?

$x_1$   
 $x_2$   
 $x_3$   
 $x_4$

.


.


.

$\vdots$   
 $x_n$

Machine  
Learning

“Dog in  
tree”





# Review: Feature engineering vs. representation learning



Dumb feature  
extraction

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ \vdots \\ x_n \end{bmatrix}$$


Representation  
Learning

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \\ \vdots \\ y_m \end{bmatrix}$$

Machine  
Learning

“Dog in  
tree”





# Review: Typical Neural Net Characteristics


So far, Deep Learning Models have things in common:

- Input Layer: (maybe vectorized), quantitative representation
- Hidden Layer(s): Apply transformation with nonlinearity
- Output Layer: Result for classification, regression, translation, segmentation, etc.
- Models used for *supervised learning*



# Changing the Objective

Today's lecture: *unsupervised learning* with neural network.





# Properties of Representations


- Smoothness
  - Small change to input  $\rightarrow$  small change in representation
- Compactness ... or overcompleteness
  - Compress input / remove redundancy
  - Multiple possible representations, chosen via other considerations (like sparsity)
- Sparsity
  - For any given input, most features are zero
- Disentangle factors of variation
  - Feature tend to vary independently of each other
- Abstraction / Invariance
  - Recognizing high-level properties shared between superficially dissimilar inputs



# Relationship between Principal Component Analysis and Autoencoder

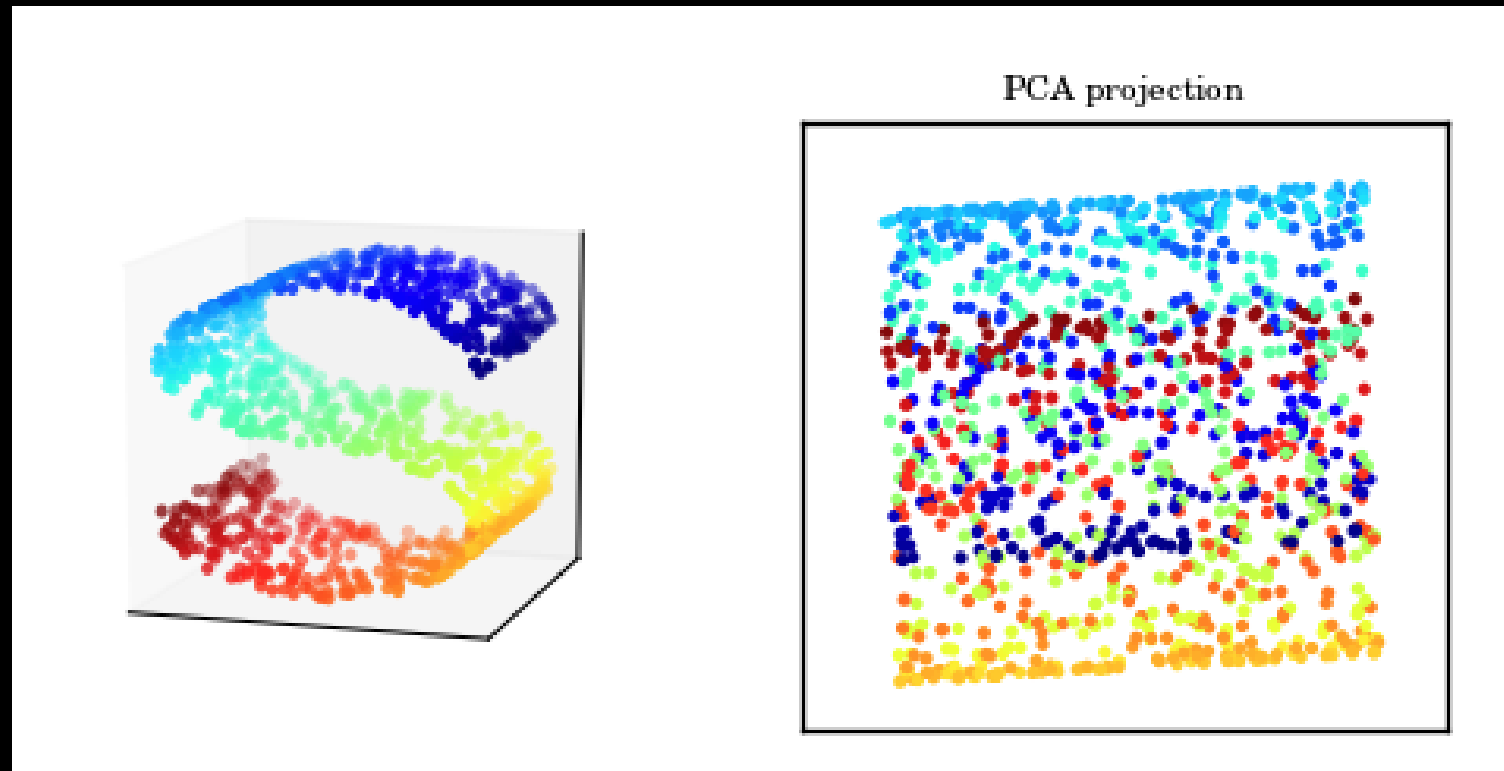
- PCA takes a N-dimensional data and represented using a much lower dimensional code.
- This happens when data lies near a linear manifold in the high dimensional space.
  - Manifold  $\rightarrow$  is a topological space that all along that space there is very high probability of our data belong to that area.






# Visualizing Manifolds

Extract 2D manifold of data which exist in 3D:






# Relationship between Principal Component Analysis and Autoencoder


- PCA takes a N-dimensional data and represented using a much lower dimensional code.
- This happens when data lies near a linear manifold in the high dimensional space.
  - Manifold  $\rightarrow$  is a topological space that all along that space there is very high probability of our data belong to that area.
- You can do this efficiently by using PCA or you can do it ineffecently by using linear Autoencoder.



# A picture of PCA with $N=2$ and $M=1$

- This takes 2-dimensional data and finds the 1 orthogonal directions in which the data have the most variance.
- The red point is represented by the projection of the red point onto the line. The blue point has an error equal to the distance from the red point to the line.





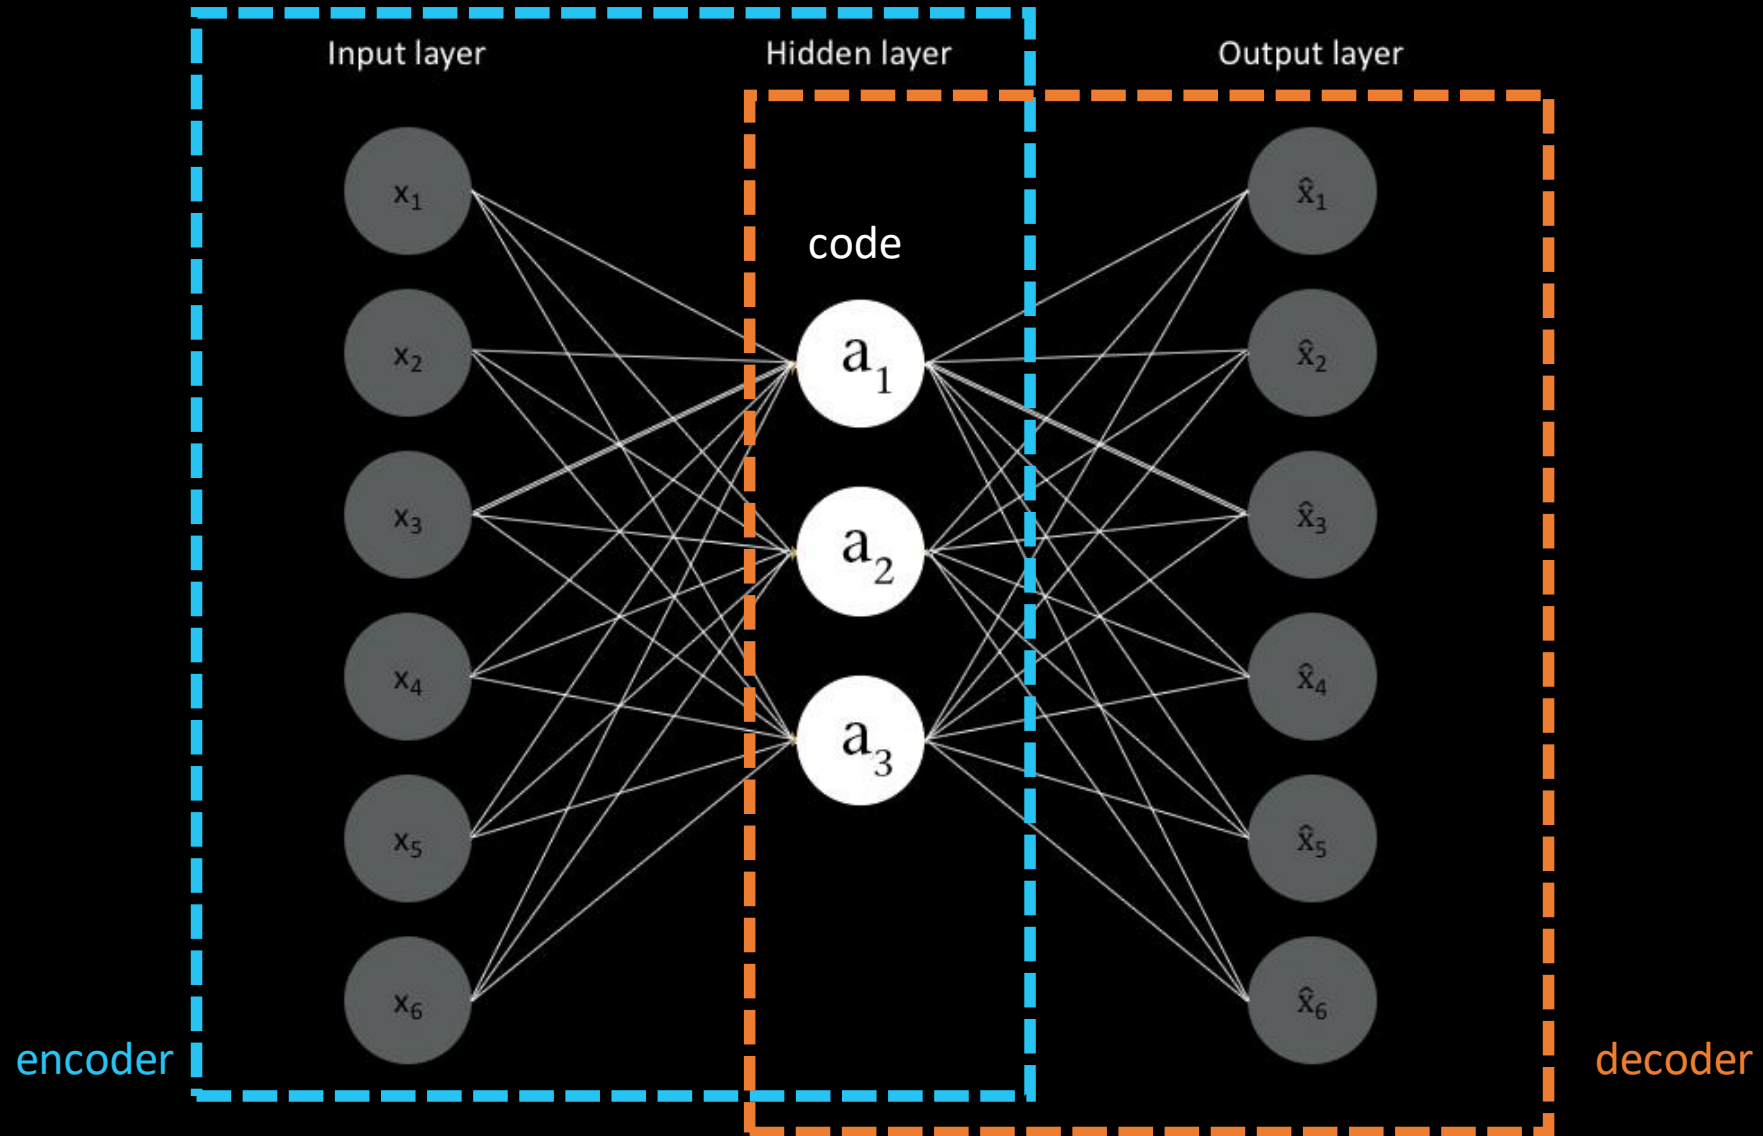
# An inefficient Autoencoder to implement PCA


Autoencoders are neural networks that are trained to copy their inputs to their outputs. Therefore, the objective function is  $\min |x - \hat{x}|$ .

- If hidden and output layers are linear it will learn hidden units that are a linear function of the data and minimize the squared reconstruction error.
  - This is exactly what PCA does.



# Visualization of a simple Autoencoder






# Code Demo of a simple Autoencoder


Let's implement a simple linear autoencoder and compare it to PCA.






# Deep Autoencoder

We learn how to implement a linear auto encoder, let's try to make it nonlinear by adding activation function and more hidden layers.




# Visualizing Deep Autoencoder





# Demo Deep Autoencoder


Let's look at the implementation.



# Caveats and Danger

Unless careful, autoencoders will not learn meaningful representations.


- Reconstruction loss: *indifferent to latent space (code)* characteristic. (not true for PCA)
- Higher representational power gives flexibility for suboptimal encoding.
- Pathological case: hidden layer is only one dimension, learns index mappings:  $\mathbf{x}^{(i)} \rightarrow i \rightarrow \mathbf{x}^{(i)}$ 
  - Not very realistic, but completely possible.




# Variational Autoencoder

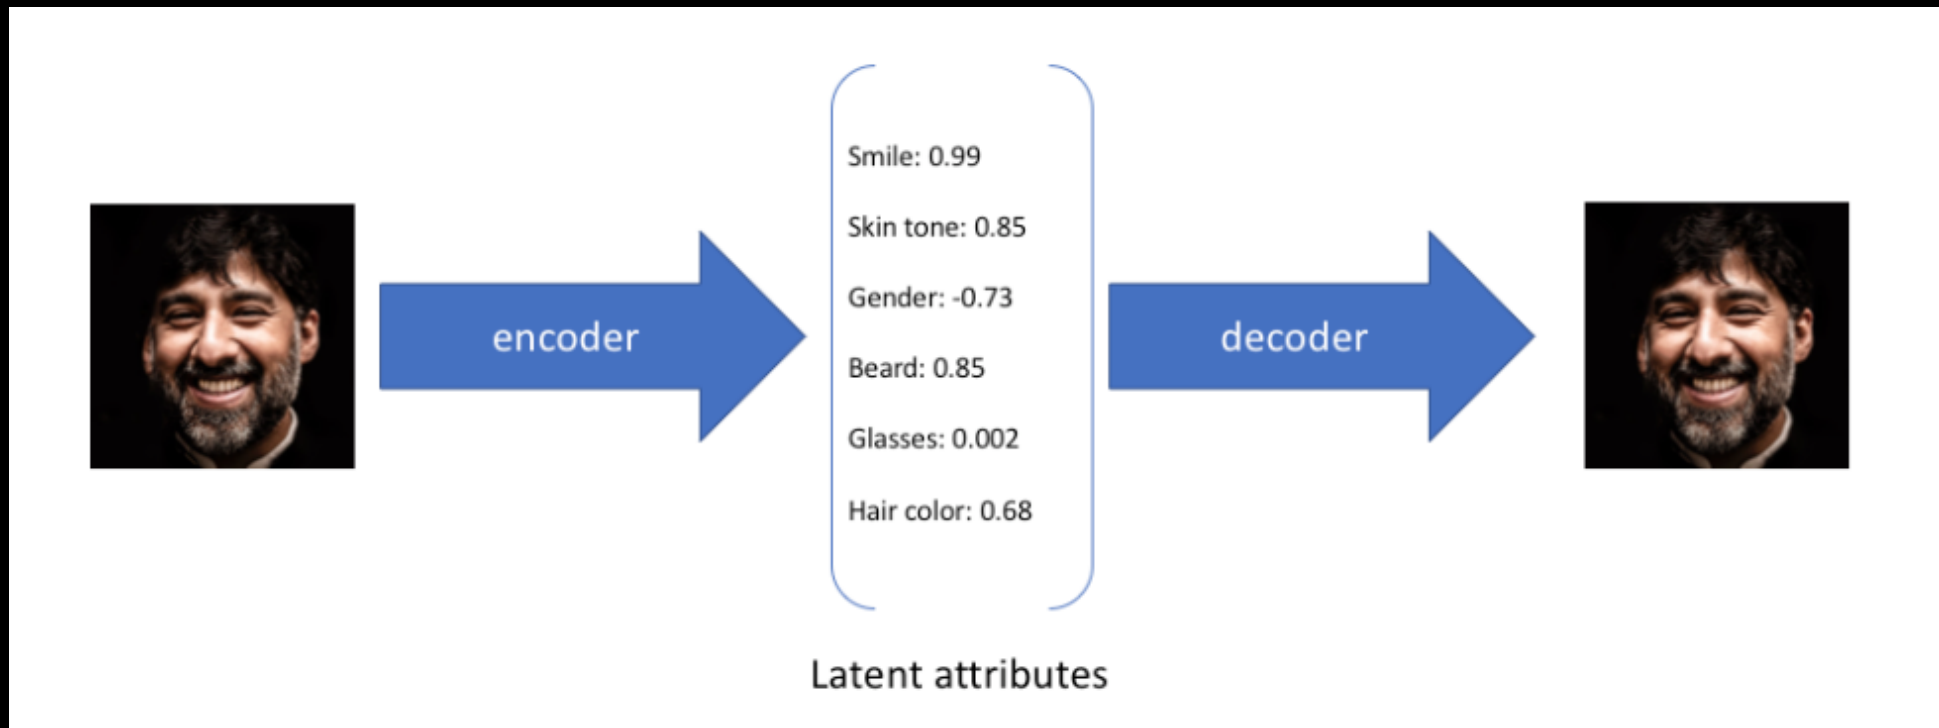
Idea: Latent space explicitly encodes distribution! Typically made to encode unit Gaussian.

- Flow: Input  $\rightarrow$  encode to statistic vectors  $\rightarrow$  sample a latent vector  $\rightarrow$  decode for reconstruction
- Loss: Reconstruction + KL Divergence





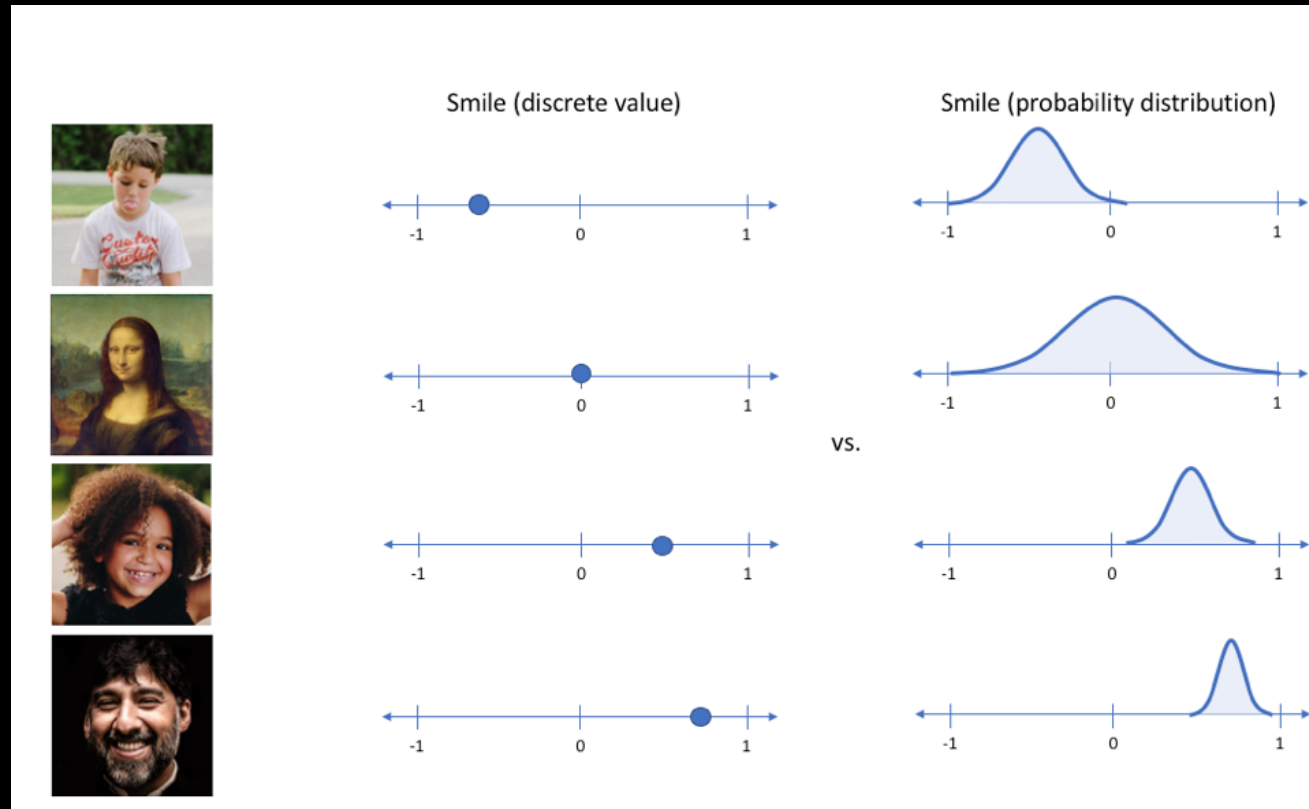
# Autoencoder vs. Variational Autoencoder









# Autoencoder vs. Variational Autoencoder



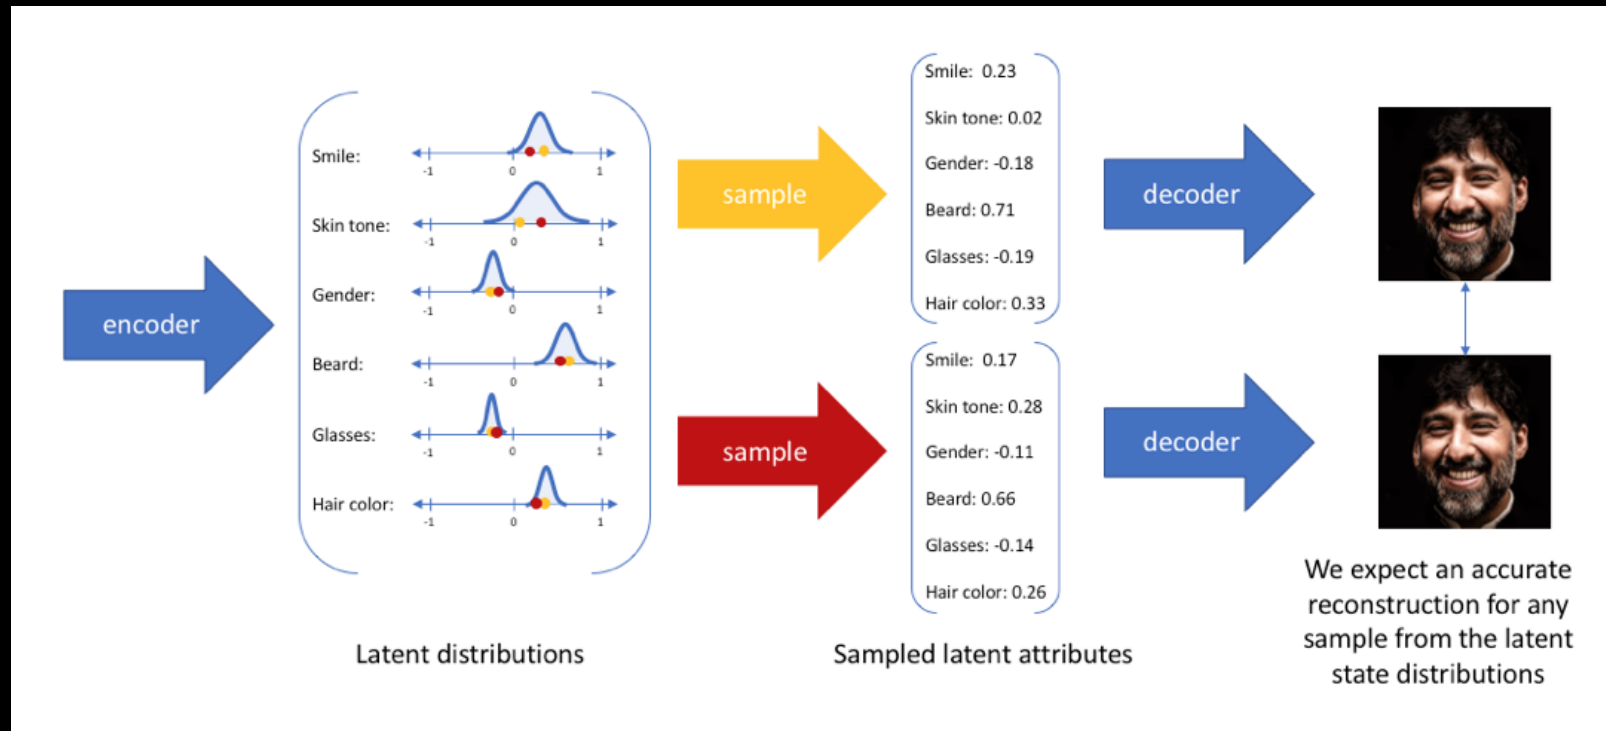



# Autoencoder vs. Variational Autoencoder





# Autoencoder vs. Variational Autoencoder





# KL Divergence

- Information Gain

$$I = -\log p(x), \quad x \in event$$

- Entropy


$$H = - \sum P(x) \log P(x)$$

- KL Divergence

- let's have two distribution  $p$  and  $q$ , KL divergence calculates dissimilarity between these two distribution

$$KL(p||q) = - \sum p(x) \log q(x) + \sum p(x) \log p(x)$$

$$KL > 0 \quad KL(p||q) \neq KL(q||p)$$



# Variational Inference


Problem definition

- Observation Data:  $x = \{x_1, x_2, \dots, x_n\}$
- Hidden Variable:  $z = \{z_1, z_2, \dots, z_n\}$

$$p(z|x) = \frac{p(z, x)}{p(x)} = \frac{p(x|z)p(z)}{\int p(x|z)p(z)dz}$$

Unfortunately, computing  $p(x)$  is quite difficult.






# Variational Inference

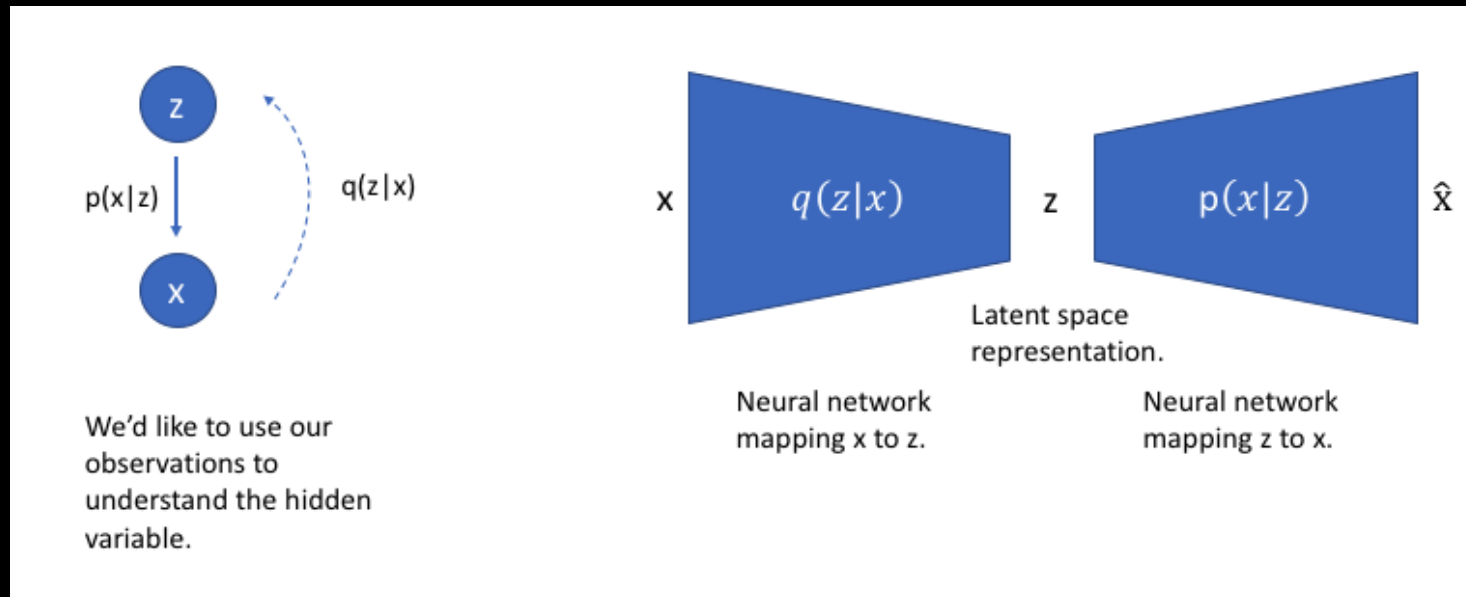
- To solve the problem, we are going to apply variational inference.
  - Let's approximate  $p(z|x)$  by another distribution  $q(z|x)$  which we'll define such that it has a tractable.
- Our goal is to make  $q(z|x)$  similar to  $p(z|x)$ 
  - $\min KL(q(z|x) || p(z|x))$
  - We can minimize the above expression by maximizing the following:
$$E_{q(z|x)} \log p(x|z) - KL (q(z|x) || p(z))$$
    - The first term represents the reconstruction likelihood.
    - The second term ensures that our learned distribution  $q$  is similar to the true prior distribution  $p$ .






# Variational Autoencoder


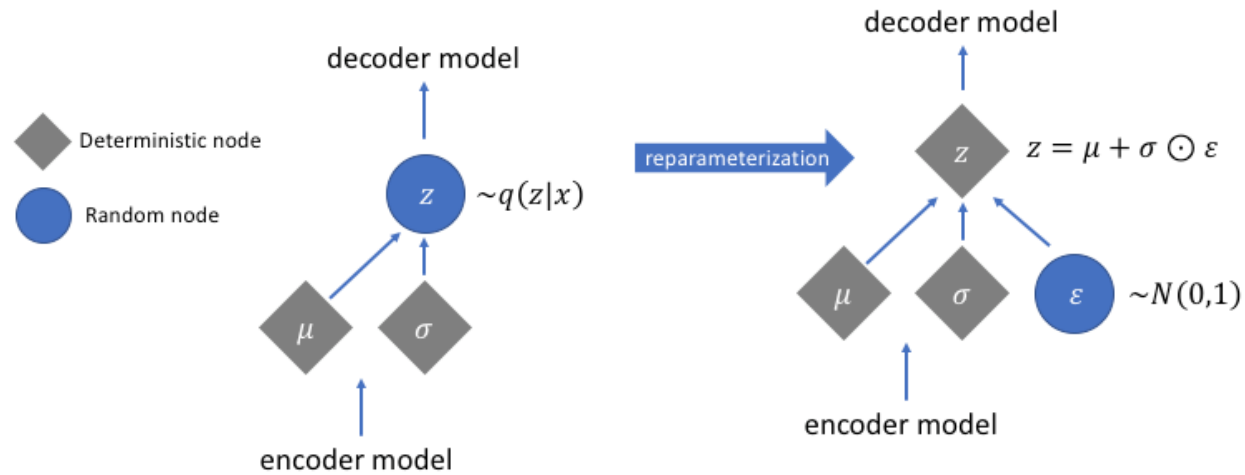
Let's revisit our network




$$\text{Loss} = J(\theta) = L(x, \hat{x}) + \sum_j KL(q_j(z|x) || p(z))$$




# Representation Trick





# Code Demo

Let's implement a variational autoencoder.



# Questions

